11.08.2025 – Task

Delta Lake Operations On Azure Databricks

1. Create a Delta Table

We start by loading a CSV file into a Spark DataFrame and saving it as a Delta table.

```
# Step 1: Load CSV file into a DataFrame

file_path = "/Volumes/workspace/default/delta/export.csv"

df = spark.read.format("csv") \

    .option("header", "true") \      # Use first row as header

    .option("inferSchema", "true") \ # Automatically detect data types

    .load(file_path)

df.show(5)  # Show first 5 rows of data
```

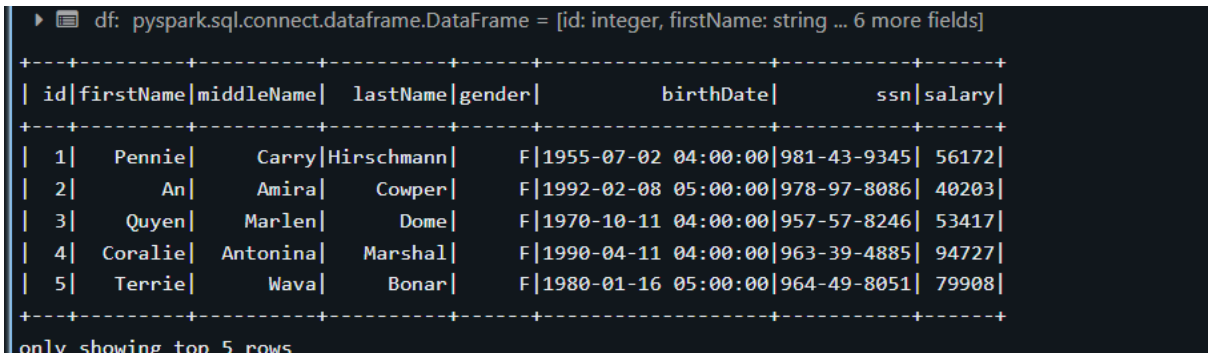

```
df: pyspark.sql.connect.dataframe.DataFrame = [id: integer, firstName: string ... 6 more fields]
+---+---------+----------+----------+------+-------------------+-----------+------+
| id|firstName|middleName|  lastName|gender|          birthDate|        ssn|salary|
+---+---------+----------+----------+------+-------------------+-----------+------+
|  1|   Pennie|     Carry|Hirschmann|     F|1955-07-02 04:00:00|981-43-9345| 56172|
|  2|       An|     Amira|    Cowper|     F|1992-02-08 05:00:00|978-97-8086| 40203|
|  3|    Quyen|    Marlen|      Dome|     F|1970-10-11 04:00:00|957-57-8246| 53417|
|  4|  Coralie|  Antonina|   Marshal|     F|1990-04-11 04:00:00|963-39-4885| 94727|
|  5|   Terrie|      Wava|     Bonar|     F|1980-01-16 05:00:00|964-49-8051| 79908|
+---+---------+----------+----------+------+-------------------+-----------+------+
only showing top 5 rows
```

```
# Step 2: Save DataFrame as a Delta table named 'employees'

df.write.format("delta") \

    .mode("overwrite") \            # Overwrite if table exists

    .saveAsTable("employees")       # Save as managed Delta table
```

---

2. Upsert (Merge) Data into the Delta Table

We create new data to update existing rows or insert new ones, then perform a merge (upsert) operation.

from delta.tables import DeltaTable

```
# Prepare new data for upsert
new_data = [
    (3, "Quyen", "Marlen", "Dome", "F", "1970-10-11 04:00:00", "957-57-8246", 55000),  # Update salary for id=3
    (6, "John", "M", "Doe", "M", "1985-05-10 04:00:00", "123-45-6789", 60000)          # New employee with id=6
]
columns = ["id", "firstName", "middleName", "lastName", "gender", "birthDate", "ssn", "salary"]
new_df = spark.createDataFrame(new_data, columns)


# Load Delta table and perform merge
deltaTable = DeltaTable.forName(spark, "employees")
deltaTable.alias("old").merge(
    new_df.alias("new"),
    "old.id = new.id"  # Match records on 'id'
).whenMatchedUpdate(set =
  {
    "firstName": "new.firstName",
    "middleName": "new.middleName",
    "lastName": "new.lastName",
    "gender": "new.gender",
```

```
    "birthDate": "new.birthDate",
    "ssn": "new.ssn",
    "salary": "new.salary"
  }
).whenNotMatchedInsert(values =
  {
    "id": "new.id",
    "firstName": "new.firstName",
    "middleName": "new.middleName",
    "lastName": "new.lastName",
    "gender": "new.gender",
    "birthDate": "new.birthDate",
    "ssn": "new.ssn",
    "salary": "new.salary"
  }
).execute()
```

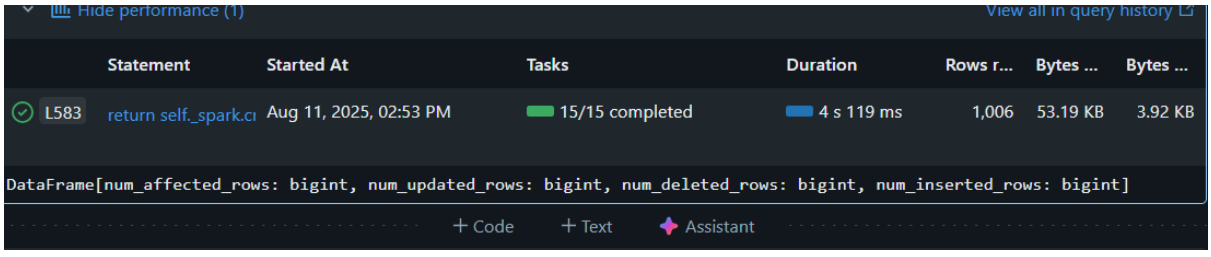

---

3. Read Data from the Delta Table

To verify the upsert operation, we read and display all data from the Delta table.

spark.sql("SELECT * FROM employees ORDER BY id").show()

id	firstName	middleName	lastName	gender	birthDate	ssn	salary
1	Pennie	Carry	Hirschmann	F	1955-07-02 04:00:00	981-43-9345	56172
2	An	Amira	Cowper	F	1992-02-08 05:00:00	978-97-8086	40203
3	Quyen	Marlen	Dome	F	1970-10-11 04:00:00	957-57-8246	55000
4	Coralie	Antonina	Marshal	F	1990-04-11 04:00:00	963-39-4885	94727
5	Terrie	Wava	Bonar	F	1980-01-16 05:00:00	964-49-8051	79908
6	John	M	Doe	M	1985-05-10 04:00:00	123-45-6789	60000
7	Geri	Tambra	Mosby	F	1970-12-19 05:00:00	968-16-4020	38195
8	Patria	Nancy	Arstall	F	1985-01-02 05:00:00	984-76-3770	102053
9	Terese	Alfredia	Tocque	F	1967-11-17 05:00:00	967-48-7309	91294
10	Wava	Lyndsey	Jeandon	F	1963-12-30 05:00:00	997-82-2946	56521
11	Sophie	Emerita	Hearn	F	1979-09-17 04:00:00	977-66-4483	90920
12	Jodie	Tabetha	Laneham	F	1959-01-31 05:00:00	923-24-9769	90634
13	Marietta	Mandi	Yansons	F	1974-02-19 04:00:00	900-34-8083	93162
14	Caridad	Maire	Snelle	F	1960-09-26 04:00:00	992-11-7062	38859
15	Yasmine	Meg	Edworthye	F	1960-01-29 05:00:00	922-12-9862	76220
16	Chan	Jani	Hartas	F	1986-12-05 05:00:00	995-51-3115	75050
17	Evangeline	Wanetta	Casserley	F	1961-09-29 04:00:00	926-61-3526	62814
18	Elnora	Kecia	Lipman	F	1980-02-14 05:00:00	950-23-9739	71350

---

4. Display the Table History

Delta Lake maintains a transaction log. We can display the full history of all operations on the table.

history_df = deltaTable.history()

history_df.show(truncate=False)